

Third Year Project  
Making Git Accessible For Children

Samuel Da Costa  
Supervisor: Caroline Jay

April 7, 2019

# Chapter 1

## Introduction

In England and Wales, the percentage of students choosing to study a GCSE in computer science rose by 50% from 2015 to 2016 and continued to rise between 2016 and 2017 [13]. This is according to official statistics published by Ofqual.

Combined with government aims to introduce simple programming to children as young as five years old, this suggests that more and more young people are beginning to produce code.

This project does not aim to convince the reader toddlers are in need of a source control system, however for students at GCSE age, and for ALevel students, both of whom are completing projects involving non-trivial amounts of code, the need for comprehensive backups and version control become more apparent.

In the field of computer science, this is a solved problem. The *Git* source control solves all these problems. Providing means of easily keeping backups, keeping track of versions and providing the ability to easily roll back to an old version.

Of course, using Git comes with its own challenges. To use it effectively requires practice and an understanding of its low-level commands and can be daunting at first.

This project aims to strike a balance between the complexity of the system and the needs of younger users. To take advantage of such a rich version control system and to provide an interface to it which can be understood and used even by children.

The emphasis of this interface is on the concept of *version control* rather than Git specifically. Abstracting and automating the internals of Git to allow the user to concentrate on the higher level concepts of, “rolling back,” and, “saving a new version”.

### 1.0.1 Background on Git

This section will briefly describe some of the main advantages of Git and introduce some relevant terminology.

As previously mentioned Git allows for easy backups. In practice, this means the existence of a remote server which also holds a copy of the user’s code. When the user wishes to back up their code. They create what is known as a `commit` object. This contains a subset of the files in the directory that the user

has edited since the last commit, as well as metadata such as the date, author e.t.c. Importantly, the metadata contains a pointer to the last commit that was created. This then forms a graph of the project history.

At any point in the future, the user now has the ability to reset the project to the state stored in any `commit`. These objects can also be pushed to the remote server, this is the way Git handles backups.

Sharing of projects across devices now becomes easy. The user on a different device makes a copy of the project from the remote server (this is known as *cloning*). They are then free to create their own commits and to push them to the remote server, and to pull down commits created by others.

This leads on to the fact Git is difficult to learn. The following is simply the author's opinion and should not be interpreted as a statement of concrete fact.

This being said, it seems clear that Git's main complexity comes from the fact that any piece of code could be in one of many places. It could exist in a stash, in the working directory, in the staging area, in a commit on the remote e.t.c.

Users are required to know about all these different areas and how they interact. When they inevitably make a mistake, usually this results in an error which is very difficult to read and interpret.

With this in mind, the first goal of this project becomes clear. To abstract away as many of these different places as possible. In fact, we will see that the project will present the user with only two areas. The remote version, and their local version.

## 1.0.2 Project Goals

On the outset of the project, there were two main objectives, from which all other aims extend. These are as follows.

**Objective 1:** To produce a user interface which should be useable by anybody over the age of thirteen years old. This interface should enable interaction with an underlying Git repository at a very high level of abstraction. *Note:* The GUI is not aiming to *teach* the user about Git, or even about version control. The idea, is that given a brief introduction to the concept of version control, that the user should be able to easily use the system.

What follows is the list of use cases the project originally aimed to fulfil. The one use case highlighted in red was later judged to add needless complexity and was not implemented.

- Ability to login on any machine with the software installed and see a list of projects.
- Ability to open any of these projects locally just with a click.
- New commits should appear on other machines logged into the same user account without the need for the user to perform manual steps.
- Ability to jump back and forwards through the project history.
- **Ability to choose which files to add to a commit**

**Objective 2** The second objective is less tangible and is best described as follows. To produce a system that is scalable, well tested and robust enough for

use by inexperienced users. In practice, this last part implies that the system should be able to automatically detect and correct issues with the underlying Git repository which would usually need to be manually corrected.

Scalability is an important part of this project and warrants further discussion. Consider the case that the project is wildly successful and has many users beginning to use it. This means many repositories needing to be stored on remote servers. One of the key aims of the project is to build a system that allows more remote machines to be added to the pool of available resources and that the system automatically begins using these multiple servers.

## Chapter 2

# Technical Design

This chapter aims to describe and explain the system’s overall architecture. To explore potential performance bottlenecks and their solutions, as well as discussing the roles of the individual components that make up the final solution. This chapter does not discuss the implementation details, but rather give a high-level view of the *why* rather than the *how*. It is relevant to mention that the system as a whole has been informally dubbed ‘littlegit’, a name used throughout this report.

### 2.0.1 Other work in this area

So-called *git guis* have existed almost as long as git itself has existed. Notable, git itself comes with two visual tools, **gitk** and **git-gui**. The former is used mainly for viewing the commit history, and the latter’s main function is to create commits [9].

However, these tools provide very little functionality beyond a graphical way of invoking git commands. There is very little automation, and a reasonable knowledge of Git is still very much required. Furthermore, they provide no help with generation of SSH keys, or the resolving of conflicts/ SSH keys are the most convenient means of authentication with Git. Please see Appendix A for an overview of how Git handles authentication through SSH.

At the time of writing the most comparable system to that which this project aims to create belongs to Atlassian. The user can manage remote repositories, clone them to their local machine, view history and interact with the repository in man complex ways through the desktop application *Source Tree*. The most noticeable feature here is the management of remote repositories from within the same application.

However, here too is where the automation stops. There is little to no help with the resolving of conflicts, SSH key setup is still a manual process, and most importantly, errors from git are often directly passed onto the user with no attempt to resolve the issue for them. Figure ?? is an example of this.

The project, therefore, aims to build upon the positive features of Source Tree and the Atlassian ecosystem, adding more abstraction and automation in the hopes of achieving a much simpler user experience.




Figure 2.1: A high level representation of the architecture of the system.

## 2.0.2 Technical Overview

First, we present a very high-level description of the system as a whole, before exploring its components and their roles.

Figure 2.1 shows an overview of the system. The three main components shown run on distinct machines communicating over the internet. The GUI is the graphical user interface that runs on the user’s machine, and what they use to manipulate their repositories.

The reader may recall that to take advantage of Git fully as a source control system the user must define a *remote repository*. A repository hosted on a remote server which remains in sync with the local copy on the user’s machine. The Git servers in Figure 2.1 are the machines that host these remote repositories.

The final component is the main server. It does *not* host remote repositories but manages the system as a whole. The main server keeps track of users and their repositories, as well as ensuring that the correct users have access to the correct repositories.

### Git Servers and the Main server

The user may ask whether this infrastructure is needed. There already exist many hosting services for Git repositories. GitHub, Gitlab and Bitbucket to name a few. However, the issue with using these remains to be our goal of automating as much as possible. To be able to set up SSH keys for the user we must have complete control over the backend infrastructure.

Furthermore, this control allows for easy integration of new features in future. For example, we could introduce a new type of user such as a teacher




Figure 2.2: A representation of the interaction between the GUI and the littlegit-core

who could have access to multiple repositories. To introduce such a role while relying on the stability of a third party service would be much more difficult.

Figure 2.1 shows multiple Git servers; this demands some explanation. A reasonable question to ask is why is one not sufficient? The second of the two objectives for this project is to produce a scalable system. We must ask ourselves what happens if we start to get more users, creating more repositories. One server soon runs out of space. Moreover, our system is likely to slow when one Git server alone has multiple other machines communicating with it simultaneously.

For this reason, we spread the load of Git repositories over multiple machines. A future goal of this project is to have machines hosted in different regions around the world, and assign their repositories based on the location of users to make the system as efficient as possible.

Of course, having multiple servers hosting Git repositories presents its own set of challenges. For a user's machine to be able to communicate with a Git server, we must register its public SSH key on all the Git servers that host the user's repositories. One of the responsibilities of the main server is precisely this. To manage SSH keys, ensuring the right ones exist on the correct machines, which becomes more complicated when we realise that an SSH key is associated to a machine, not a user and that a user may work on multiple machines necessitating multiple keys.

The management of SSH keys is only one role assigned to the main server; it also creates Git repositories on the Git servers and manages users, keeps track of their repositories and handles all user authentication.

## GUI

The GUI is the desktop application the end user interacts with and uses. Its primary responsibilities are to allow the user to manipulate the Git repository and visually display its current state.

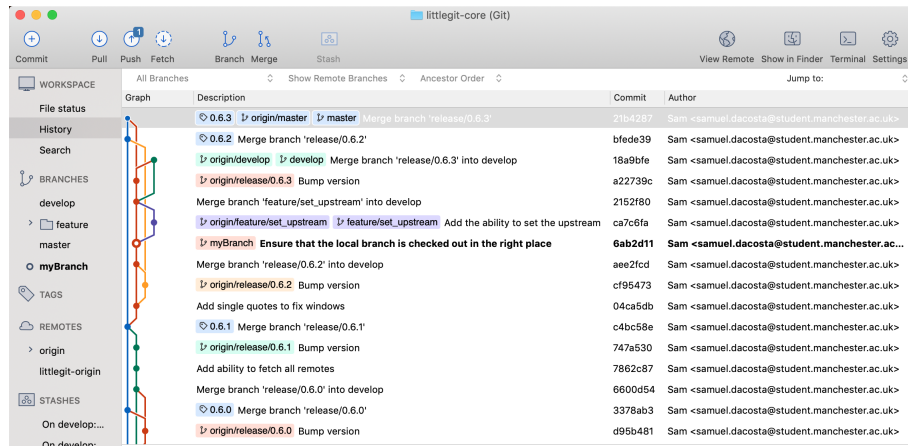


Figure 2.3: A screenshot of the main interface of the application SourceTree.

### Low-level git library

At this stage, it's prudent to mention the way both the GUI the main server interact with Git itself. This interaction happens through the medium of a library called the **littlegit-core** written for this system. It serves as a wrapper around the Git command line allowing interaction from a high-level language (discussed further in chapter 3) with the Git binary. The littlegit-core is responsible for executing git commands and parsing the output into useful high-level objects.

However, the reader should note that this layer performs none of the automation of Git which the system aims to achieve. This layer executes git commands then parses and returns the resulting output as Figure 2.2.

Both the main server and GUI include the littlegit-core library as a dependency, allowing them to interact with Git repositories without interacting directly with the Git binary.

We discuss the GUI's design and features in the following section.

### 2.0.3 UI Design

Early in the project, the interfaces of other Git GUIs were examined, exploring their strengths, weaknesses and most importantly, their common traits. The findings of this research are discussed here as well as a discussion of how the results impacted the design of this project's user interface.

The three main applications researched are the application SourceTree [5] (mentioned earlier), GitKraken [6] and GitK [10]. Figures 2.3, 2.4 and 2.5 show a screenshots of their main interfaces.

The first thing to note in looking at the interfaces is the clear common element being the Git graph. Though each has a different style and displays the information differently, the graphs are all large clear features of each interface. The decision was made to emulate the overall style in littlegit. The vertical graph, new commits on top with older commits in descending order below them. The reasoning here is that the Git history portrayed in this way conveys all needed information, but it is easy to present a simplified version (as was needed for littlegit) by simply omitting some of the information surrounding it.






Figure 2.4: A screenshot of the main interface of the application GitKraken.

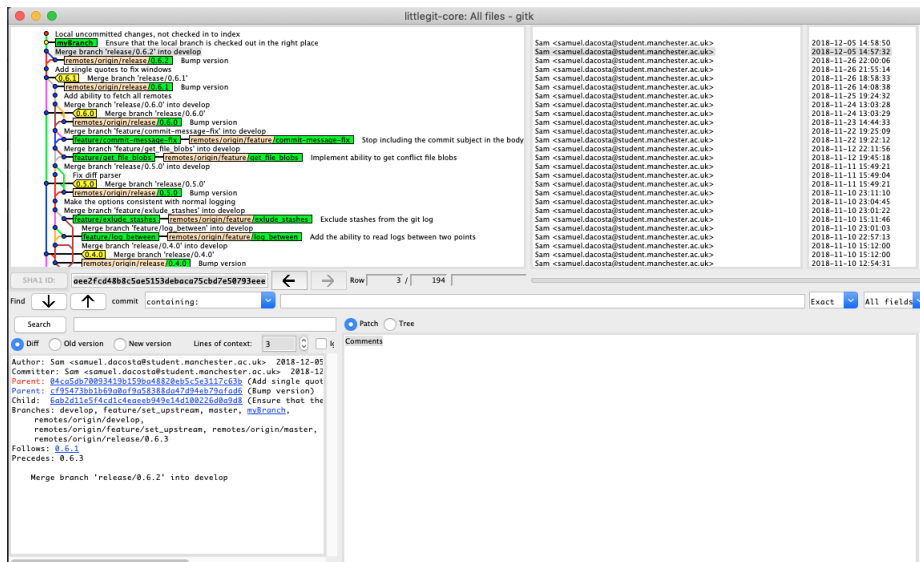


Figure 2.5: A screenshot of the main interface of the application GitK.




Figure 2.6: An early design for the Git graph using horizontal scrolling.

Though other options were considered, Figure 2.6 shows an early design for an alternative graph. The idea was to emulate a timeline to help visualise the repository as being a series of snapshots of the code, allowing the user to navigate backwards and forwards along it. This idea however soon proved impractical, firstly visually representing branches in this model becomes messy and complicated, especially for more complex repositories. Furthermore, according to user testing [12], users dislike horizontal scrolling.

Returning to our discussion of the other applications, the graphs of both SourceTree and GitK are relatively cramped. A large amount of information in very close proximity, this is especially true of GitK. On the other hand, GitKraken takes much more space to display the same information, a feature it was important for littlegit to emulate to ensure the data is clear and easy to understand.

The next point to make is that all three applications have lots of buttons and menus. By looking through all the buttons and options, it became clear that in a system that aims to automate as much as possible. The vast majority of the options were unneeded. For example, SourceTree has three separate buttons for pushing, pulling and fetching. The question littlegit asks is why? If there is something to push to the server, why not just do it on the users' behalf, and if there are no new changes on the server to be pulled, then why present the user with the option? This led to the decision of having a popup dialogue appear when the user's local version is behind the server, giving them the option to update without the need to clutter the primary interface with the option.

Since this automation is taking place, it seems prudent to explicitly describe the Git workflow littlegit aims to present. The workflow is built around the commit, creating commits should drive everything else. Creating a commit causes changes to be pushed to the server, and the arrival of new commits on the server should cause the application to update the repository accordingly. Branches in this workflow are much less significant and are only used by littlegit internally when needed. With this in mind, looking back at SourceTree especially, it is very non-obvious how to create a new commit, something imperative to avoid in littlegit.

Hence, creating new commits should be the primary operation the user sees when using the application. Furthermore, by looking at GitKraken and SourceTree in particular, the reader may also notice that many of the buttons on their left toolbars are merely the names of branches. In regular use clicking on them allows the user to check out these branches. Both applications also allow for using the graph to navigate between branches, allowing the user to click on the branch they wish to check out. With littlegit we aimed to expand this functionality, allowing the user to check out any commit in the history and handling this with branches internally without exposing this to the user.

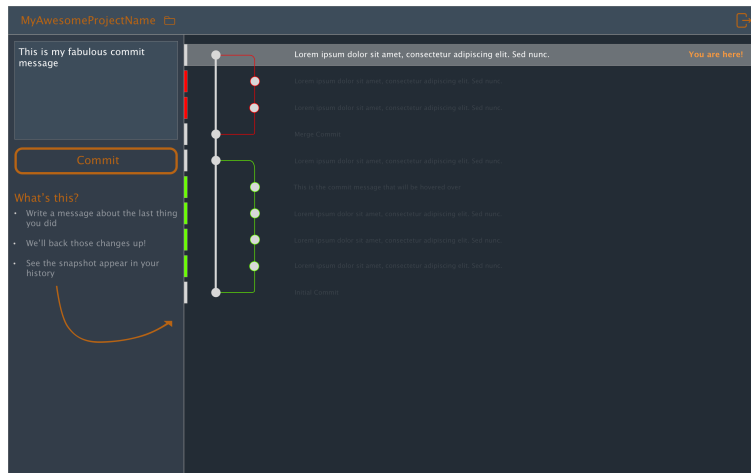


Figure 2.7: The main screen of littlegit

## 2.0.4 Designing for teenagers

There has been much research conducted on the topic of designing user interfaces for teenagers. Unfortunately, the orientation of much of that research was towards websites and mobile applications rather than desktop apps. A result of this is usually much discussion of the importance of social media integration and good mobile compatibility, neither of which apply to a desktop application such as littlegit.

However, many of the conclusions are useful, particularly a study by the Nielsen Norman Group concerning designing for teenagers [11].

One of the biggest influencers of this study on the design of littlegit is the conclusion that teenagers hate childish content, seeing it as patronising. For this reason, littlegit refrains from any unneeded graphics and multimedia. Where an explanation of functions are required, they are short and do not talk down to users.

Furthermore, the study suggests that the strategy mentioned earlier of reducing clutter on the main screen is the correct approach, including refraining from using flashy animations and graphics.

Lastly, a point continually emphasised by the study is that of fast loading speeds. Though this refers specifically to websites, the study highlights the impatience of teenagers while using the software. The second objective of the project; therefore, of a scalable architecture which allows the application to work at speed is doubly important when considering the target audience.

## 2.0.5 Final design before implementation

The development of the software followed an iterative approach. This meant continually seeking feedback from others as the project continued. However, it was essential to have an initial design to aim for, taking into account the research above.

To do this the design software Sketch [7] was used to produce initial designs. These designs are included in Figures 2.7 and 2.8.

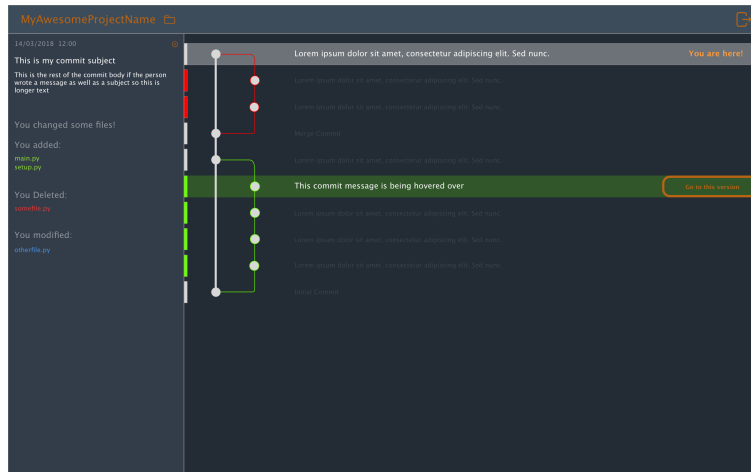


Figure 2.8: The main screen of littlegit after a user clicks a commit from the graph.

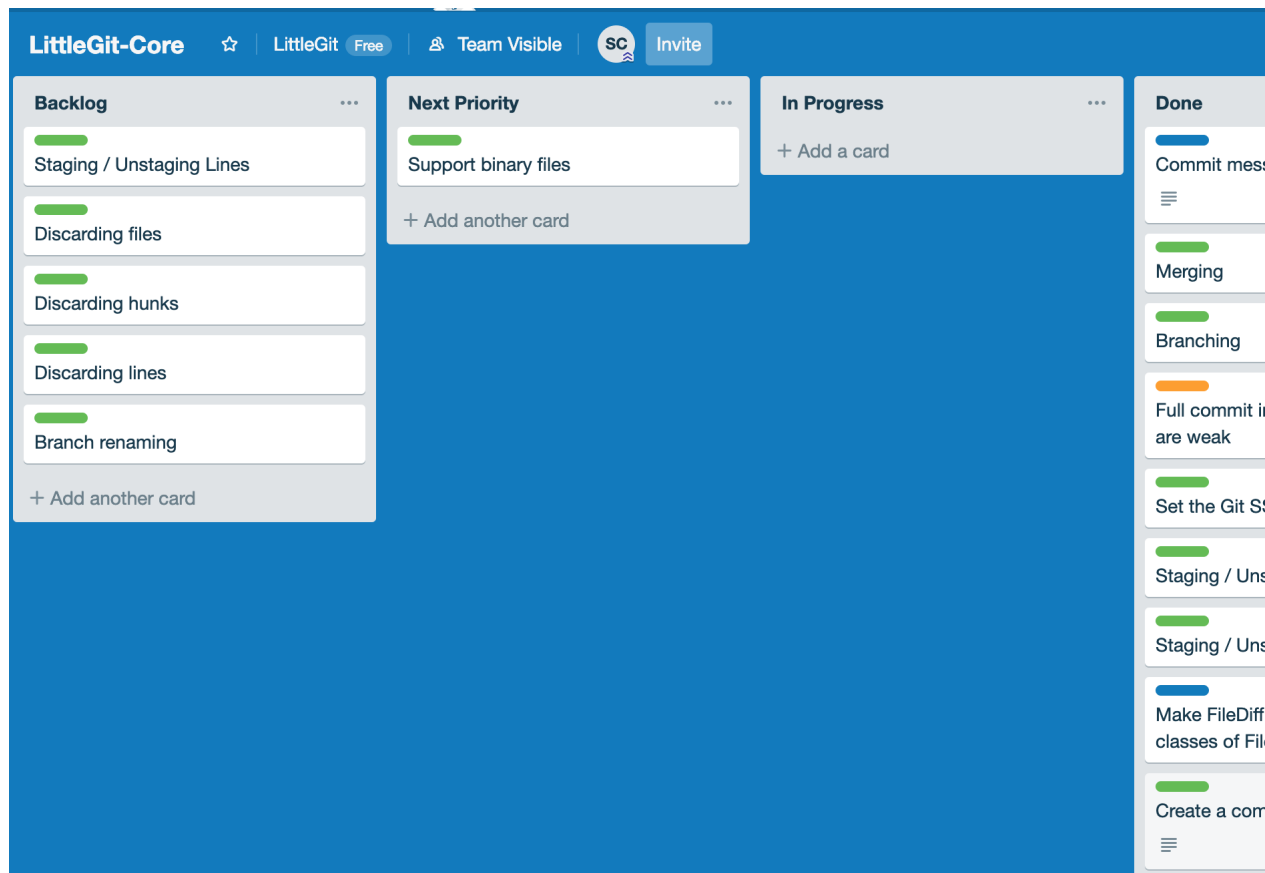


Figure 2.9: A screenshot of the Trello board for the `littlegit-core` component of the project, taken after the end of development.

In this chapter, we discuss the main development phase of the project and justify some of the main decisions made throughout its course. We will explore the implementations of the architectures discussed at a high level in Chapter 2 as well as some of the more intriguing problems encountered throughout development.

## 2.1 Planning and management of work

Trello boards were used from the very beginning to keep track of project goals and tasks. A separate board kept track of tasks relating to the server, the `littlegit-core` and the GUI. Figure 2.9 shows the current state of the board for the `littlegit-core` (though the screenshot cuts off many done tasks). The strategy used with regards to the boards was as follows. Very early in the project, a large number of tasks were created in each board and placed in the *Backlog* column. Every week a few of the highest priority tasks were moved to the next column along, the *Next Priority* column for completion that week.

The strategy was first to complete a bulk of the tasks for the `littlegit-core`,

then the server. Once both were stable, we then began work on the GUI. The approach worked reasonably well but did lead to a small amount of thoroughly tested code not being used in the final project; we will discuss the benefits and weaknesses of the approach in Section ??.

It is worth mentioning the way tasks were chosen. The approach was to take the desired feature set (discussed back in Section 1.0.2) and to break these features up into the tasks needed to achieve them.

## 2.2 Choice of language

In chapter 2 we introduced the need for a library (the *littlegit-core*) to interact with the Git binary and provide a higher level interface to it, to be used by both the main server and the GUI.

Considering the time constraints on this project, only implementing the library once was crucial. Furthermore, the time constraints would not permit writing further wrappers of the *littlegit-core* to make it compatible with other technologies. Hence the GUI and server needed to be able to import the *littlegit-core* directly.

Three main technologies were considered. They are as follows.

1. **JavaScript:** Though this technology has been a staple of web development for many years, in recent years the Electron [1] has become a common technology allowing JavaScript, HTML and CSS (all traditionally web technologies) to be used to produce cross-platform desktop applications.

Furthermore, NodeJS [3] is a technology allowing JavaScript's use in writing server code, making it a very viable choice as a core technology for *littlegit*.

However, JavaScript is a *weakly typed* language, and previous experience has indicated that maintaining a large body of weakly typed code quickly becomes very difficult. Despite the advent of tools such as *TypeScript* which is a typed language which transpiles into JavaScript, it is still no replacement for a truly typed language such as Java.

## 2.3 Implementation of components

There was an emphasis following best practices throughout this project from a software engineering point of view. Principle among these was the tenant of separation of concerns. This section details the implementation of the individual components which communicate with each other but are distinct.

## 2.4 The *littlegit-core*

As has already been discussed, this is the library that interacts directly with the Git binary. The primary challenge involved in producing this component was the interaction with the so-called git-plumbing.

Git commands are loosely grouped into two categories (though the distinction is not a hard one). The first is git-porcelain. Porcelain commands are

those with which most Git users are likely to have interacted. A good example is `git-status`, which gives a summary of the current state of the user's local Git repository. These are commands designed for use by human users; often, their outputs are tweaked between Git versions to aid human readability. This makes them unsuited for machine use; it is impossible to reliably parse output which is slightly different in every version of Git.

The git-plumbing commands, on the other hand, are generally lower level commands to the porcelain and are often used under the hood by the porcelain itself. Their outputs are considered stable and designed for interaction with other software rather than humans. While making them a far better choice for the littlegit-core to interact with, the difficulty is in the fact they are by nature lower level commands. For example, to check out a branch with git-porcelain is a straightforward one-line operation, but this conceals multiple plumbing commands under the hood. Littlegit-core endeavours to use plumbing commands wherever possible.

As has been mentioned, the project aims to follow best practices wherever possible. For this reason, and the fact that the rest of the project relies heavily on the littlegit-core being stable, a TDD testing approach was used. One way of achieving this was to *design the library for testability*. The main design decision supporting this was the decision to make the library's operations synchronous. The user's interaction with the GUI should not grind to a halt because a shell command is executing. Hence there must be some asynchronicity, but the decision was made that the GUI would be responsible for threading calls to the littlegit-core, leaving it to return its results synchronously.

Methods who return synchronously prove to be a much simpler and more reliable to test, allowing the entire library to be covered thoroughly with tests.

## Architecture

The primary driver behind the architecture of the littlegit-core is the fact it must be able to execute its commands on remote machines over an SSH connection as well as on the local machine. For example, it is used to initialise remote repositories via SSH on the Git servers.

Figure 2.10 shows the way the library is architected. The top layer is used to combine plumbing commands; for example, multiple plumbing commands are required to change branches. It makes calls to the `command runner` which provides a high-level interface to these Git commands. The `command runner` then converts these high-level invocations raw commands to be interpreted by Git, which are executed either through the local or remote runners depending on the configuration. Local for the GUI, remote on the server.

## Deployment

We keep mentioning that the library is *used* or *included* by the GUI and server. It is worth explaining what this means and describing the deployment process for the library. Figure 2.11 shows the deployment pipeline, triggered by a push to source control; all the automated tests are run remotely by Travis's continuous integration service. Importantly when the tests succeed the library is deployed to a maven repository allowing it to easily be included as a maven dependency by the server and the GUI, without the need to manually copy any files.




Figure 2.10: The internal architecture of the littlegit-core.

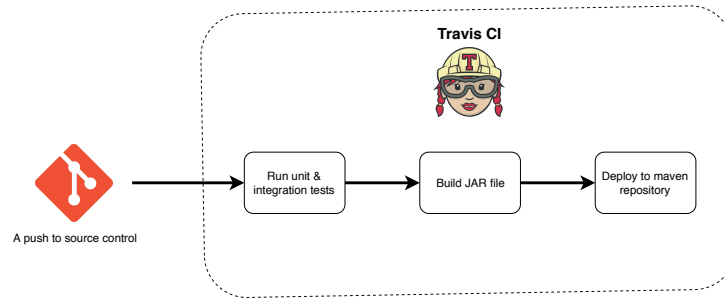


Figure 2.11: The deployment pipeline of the littlegit-core.

## 2.5 The main server

In this section, we will explore the implementation of the main server and justify the technologies used. Recall that the main server's primary function is the management of users and remote repositories. Keeping track of user details as well as information on which users have access to which repositories.

### 2.5.1 Technologies

We have already justified the decision to use the Kotlin programming language for the server. However, multiple options were available in regards to different technologies for everything from the database to the framework used to build the server itself. We discuss and justify the choices made here.

Firstly, Amazon Web Services hosts all the online infrastructure for this project. The justification for this is simple, hosting of databases and servers is expensive, and AWS offers a free tier of the services needed for students.

#### Web server framework

Let us first discuss the choice of the server framework used. The two frameworks most seriously considered are as follows.

- The Spring framework [14].



- The Jersey framework [2].

We chose Jersey for its lightweight nature, it provides an easy way of creating API endpoints but unlike Spring does not attempt to control the entire stack down to the database allowing much more flexibility. Considering the need for the server to communicate via SSH with the Git servers, having full control of the stack was vital.

The downside of Jersey and the advantage of Spring is a direct consequence. Spring provides much of the required functionality (communication with the database and authentication of users, being most significant) as standard, while these must be done manually with Jersey. In the end, the need for flexibility made Jersey the correct choice, despite the additional implementation work involved.

### Database and caching

Given the rise of non-relational databases such as CouchDB and Firebase the decision to use a traditional SQL database is no longer a trivial one. These non-relational databases have some significant advantages, one of the most enticing being their quality of scaling horizontally, unlike traditional databases which generally only scale vertically.

However, the decision was made to use a traditional database for the following reasons. The first and most important of these reasons are the need for complex querying on the data where non-relational databases do not compete with the power of traditional SQL. An example of this kind of complex query is the need to find all the Git servers which contain repositories the user either owns or has access to which do not currently hold their SSH keys.

Moreover, the data has an evident structure, a structure which does not need to change often or adapt to a rapidly changing environment (where non-relational databases come into their own). Because of this, the static, rarely changing schemes of traditional SQL databases are not a hindrance.

The scaling issue, however, must be addressed. No matter the convenience to the programmer, the use of a technology that will not scale efficiently cannot be justified when one of the objectives of the project is to produce a scalable system.

Let us consider the data that the system needs to store. For the reader with experience with relational databases, Figure 2.12 shows an entity relationship diagram for the database. In short, the data we need to store is as follows.

- User information including tokens used for authorisation and authentication.
- Information about repositories and the servers upon which they are stored.
- The SSH keys for each user.
- Information regarding which users have access to which repositories.

The scalability of the server relies on one key observation, the fact that the reading from the database is a far more common operation than writing. For example, the creation of repositories is not a common operation but checking if a user has access to a repository is very common (this will be discussed further




Figure 2.12: An entity relationship diagram representing the database.

in subsequent sections). This observation allows us to introduce caching to solve our scalability issue. By caching as many of the values as possible in a scalable, **in-memory** caching solution, we massively reduce traffic to the database. We will discuss this further in subsequent sections, for now, we will justify the use of Redis [4] as our caching solution.

Redis's feature set is vast [4], but the important ones for us are its in-memory key-value store and built-in replication. Its in-memory nature allows for very fast reads of the data. Of course, key-value pairs do not offer the benefit of flexible SQL-like queries. However, this isn't what we need at this point, if we need to do complex queries we can access the database, the cache speeds up common simple operations such as returning the details for a particular user, or checking if an authentication token is valid **without** the overhead of checking the database which is on disk.

Moreover, having a cache with built-in replication allows it to scale across multiple machines easily giving us the horizontal scaling we require. The reader may wonder why we are using a caching solution rather than making use of data structures within the server code itself to achieve the same thing. It is a very valid argument to say that such an approach would be more straightforward and easy to implement, with less to go wrong!

Let us consider the situation where the load on the server becomes large, causing requests to become slow and even to time out. Since the memory of the machine hosting the server caches the data, the only option is to get a bigger machine with more memory (vertical scaling). On the other hand, using our approach we can spin up multiple machines hosting the server to balance the load, and Redis's inbuilt replication allows us to do the same thing with the cache. Figure 2.13 shows the idea <sup>1</sup>.

## Data transfer

The decision to use the JSON data format for communication between the server and desktop client was a straightforward one. JSON is lightweight and is easily parsable by machine while also being straightforward to understand by humans (making debugging easy) [8].

<sup>1</sup>This was only tested with one machine each for the cache, server and database without a load balancer due to the cost of doing so on Amazon Web Services and other platforms, the code however fully supports the architecture described.




Figure 2.13: The overall architecture of the main server and the way it is designed to scale.




Figure 2.14: The internal architecture of the main server.

Furthermore, the decision to create a RESTful interface with the server allows it to be used by other clients in future easily. For example, if a web interface was required to enable users to view and manipulate their remote repositories, the same endpoints could be used without any additional work.

## 2.5.2 Internal server architecture

While discussing technologies, we saw the overall structure of the server technology stack. In this section, however, we discuss the way the code which interacts with these various components is architected.

As was described in Section 2.5.1, a lightweight server framework was selected for the system which does not enforce a particular architecture. The Spring framework was a loose inspiration to our architecture, resulting in the use of a *Layered Architecture*. Figure 2.14 shows the different layers of the system.

As the diagram shows, each layer abstracts the one beneath it, and each layer interacts exclusively with the one below it. The roles of the individual layers are as follows.

## Repository Layer

The repository layer is little more than an abstraction on the database and caching layers. Used by the service layer to retrieve and write information, this layer invokes raw SQL commands to communicate with the database and crucially also reads from the cache, and invalidates cached information when data is updated. For the interested reader, the caching strategy is described in more depth in Appendix C.

Furthermore, the repository layer handles the mapping of data from the database and cache to Kotlin data structures.

By the nature of this layer's direct communication with the SQL database, it is here that considerations must be made to protect against SQL injection attacks.

## Service Layer

In this layer, we have all of our application-specific business logic. For example, the process of creating a new repository consists of several steps, involving multiple writes to the database and the running of commands on remote Git servers. The logic for combining these operations exists in the service layer.

## Controller Layer

This layer defines the interface between the server's logic and the outside world. This means defining the individual API endpoints and the mappings from HTTP URLs to methods. Importantly, in this layer, we also define what kind of user can access each endpoint. A discussion of the types of users supported by the server is included in Appendix B.

By far, this is the most difficult layer to test since it how we define the interface of our API using syntax provided by the Jersey framework. For this reason, it is vital that the controller contains as little business logic as possible. In fact, for this project, each controller does nothing but invoke code in the Service layer which is far easier to test.

## 2.6 Authentication

There are two main kinds of authentication used by the system. SSH authentication is used when the desktop GUI performs Git operations interacting with the remote Git server (for example `git push` and `git fetch`), whereas calls to the API use a token-based system.

### 2.6.1 API authentication

We discuss the SSH authentication in the next section

## Appendix A

# SSH Authentication

## Appendix B

### User roles

## Appendix C

# Caching strategy

# Bibliography

- [1] Electron.
- [2] Jersey.
- [3] Nodejs.
- [4] Redis.
- [5] Atlassian.
- [6] axosoft.
- [7] Bohemian Coding. Sketch.
- [8] ereger. What is json?
- [9] git scm. Git in other environments - graphical interfaces. Accessed: 26-02-2019.
- [10] gitk.
- [11] Alita Joyce and Jakob Nielsen. Teenagers ux: Designing for teens.
- [12] Jakob Nielsen. Scrolling and scrollbars.
- [13] Ofqual. Gcse outcomes in england. Accessed: 26-02-2019.
- [14] Pivotal. The spring framework.